

# Lecture 26 — Profilers, Profiler Guided Optimization

Patrick Lam & Jeff Zarnett

`patrick.lam@uwaterloo.ca, jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering  
University of Waterloo

December 5, 2020

# Part I

## Profiling Tools

# No-One Expects the Profiling Tools!



**AMONGST** our profiling tools are such diverse elements AS...

- Solaris Studio Performance Analyzer
- VTune (Intel)
- CodeAnalyst (AMD)

Plus a few more we'll consider in some more detail.

Intrumentation-based tool.

System-wide.

Meant to be used on production systems.



(Typical instrumentation can have a slowdown of 100x (Valgrind).)

Design goals:

- 1 No overhead when not in use;
- 2 Guarantee safety—must not crash  
(strict limits on expressiveness of probes).

How does DTrace achieve 0 overhead?

- only when activated, dynamically rewrites code by placing a branch to instrumentation code.

Uninstrumented: runs as if nothing changed.

Most instrumentation: at function entry or exit points.

You can also instrument kernel functions, locking, instrument-based on other events.

Can express sampling as instrumentation-based events also.

You write this:

---

```
syscall::read:entry {  
    self->t = timestamp;  
}  
  
syscall::read:return  
/self->t/ {  
    printf("%d/%d spent %d nsecs in read\n"  
        pid, tid, timestamp - self->t);  
}
```

---

`t` is a thread-local variable.

This code prints how long each call to `read` takes, along with context.

To ensure safety, DTrace limits expressiveness—no loops.

- (Hence, no infinite loops!)

Built for production environments.

Specialized for profiling JVMs,  
uses JVM hooks to analyze idle time.

Sampling-based analysis; infrequent samples  
(1–2 per minute!)



At each sample: records each thread's state,

- call stack;
- participation in system locks.

Enables WAIT to compute a “wait state”

(using expert-written rules):

what the process is currently doing or waiting on, e.g.

- disk;
- GC;
- network;
- blocked;
- etc.

You:

- run your application;
- collect data (using a script or manually); and
- upload the data to the server.

Server provides a report.

- You fix the performance problems.

Paper presents 6 case studies where WAIT identified performance problems.

Profiling: Not limited to regular compiled program code.

You can profile Python using `cProfile`; standard profiling technology.

Google's Page Speed Tool: profiling for web pages—how can you make your page faster?

- reducing number of DNS lookups;
- leveraging browser caching;
- combining images;
- plus, traditional JavaScript profiling.

I ran the command `nvprof target/release/nbody-cuda`.

==20734== Profiling application: target/release/nbody-cuda

==20734== Profiling result:

|                 | Type | Time(%) | Time     | Calls | Avg      | Min      | Max      | Name                       |
|-----------------|------|---------|----------|-------|----------|----------|----------|----------------------------|
| GPU activities: |      | 100.00% | 10.7599s | 1     | 10.7599s | 10.7599s | 10.7599s | calculate_forces           |
|                 |      | 0.00%   | 234.72us | 2     | 117.36us | 100.80us | 133.92us | [CUDA memcpy HtoD]         |
|                 |      | 0.00%   | 94.241us | 1     | 94.241us | 94.241us | 94.241us | [CUDA memcpy DtoH]         |
| API calls:      |      | 97.48%  | 10.7599s | 1     | 10.7599s | 10.7599s | 10.7599s | cuStreamSynchronize        |
|                 |      | 1.92%   | 211.87ms | 1     | 211.87ms | 211.87ms | 211.87ms | cuCtxCreate                |
|                 |      | 0.54%   | 59.648ms | 1     | 59.648ms | 59.648ms | 59.648ms | cuCtxDestroy               |
|                 |      | 0.04%   | 4.8704ms | 1     | 4.8704ms | 4.8704ms | 4.8704ms | cuModuleLoadData           |
|                 |      | 0.00%   | 404.72us | 2     | 202.36us | 194.51us | 210.21us | cuMemAlloc                 |
|                 |      | 0.00%   | 400.58us | 2     | 200.29us | 158.08us | 242.50us | cuMemcpyHtoD               |
|                 |      | 0.00%   | 299.30us | 2     | 149.65us | 121.42us | 177.88us | cuMemFree                  |
|                 |      | 0.00%   | 243.86us | 1     | 243.86us | 243.86us | 243.86us | cuMemcpyDtoH               |
|                 |      | 0.00%   | 85.000us | 1     | 85.000us | 85.000us | 85.000us | cuModuleUnload             |
|                 |      | 0.00%   | 41.356us | 1     | 41.356us | 41.356us | 41.356us | cuLaunchKernel             |
|                 |      | 0.00%   | 18.483us | 1     | 18.483us | 18.483us | 18.483us | cuStreamCreateWithPriority |
|                 |      | 0.00%   | 9.0780us | 1     | 9.0780us | 9.0780us | 9.0780us | cuStreamDestroy            |
|                 |      | 0.00%   | 2.2080us | 2     | 1.1040us | 215ns    | 1.9930us | cuDeviceGetCount           |
|                 |      | 0.00%   | 1.4600us | 1     | 1.4600us | 1.4600us | 1.4600us | cuModuleGetFunction        |
|                 |      | 0.00%   | 1.1810us | 2     | 590ns    | 214ns    | 967ns    | cuDeviceGet                |
|                 |      | 0.00%   | 929ns    | 3     | 309ns    | 230ns    | 469ns    | cuDeviceGetAttribute       |

Oh, and for comparison, here's the one where I make much better use of the GPU's capabilities (with better grid and block settings):

=22619== Profiling result:

|                 | Type | Time(%) | Time     | Calls | Avg      | Min      | Max      | Name                       |
|-----------------|------|---------|----------|-------|----------|----------|----------|----------------------------|
| GPU activities: |      | 99.92%  | 417.53ms | 1     | 417.53ms | 417.53ms | 417.53ms | calculate_forces           |
|                 |      | 0.06%   | 236.03us | 2     | 118.02us | 101.44us | 134.59us | [CUDA memcpy HtoD]         |
|                 |      | 0.02%   | 93.057us | 1     | 93.057us | 93.057us | 93.057us | [CUDA memcpy DtoH]         |
| API calls:      |      | 52.09%  | 417.54ms | 1     | 417.54ms | 417.54ms | 417.54ms | cuStreamSynchronize        |
|                 |      | 26.70%  | 214.00ms | 1     | 214.00ms | 214.00ms | 214.00ms | cuCtxCreate                |
|                 |      | 13.63%  | 109.26ms | 1     | 109.26ms | 109.26ms | 109.26ms | cuModuleLoadData           |
|                 |      | 7.42%   | 59.502ms | 1     | 59.502ms | 59.502ms | 59.502ms | cuCtxDestroy               |
|                 |      | 0.05%   | 364.08us | 2     | 182.04us | 147.65us | 216.42us | cuMemcpyHtoD               |
|                 |      | 0.04%   | 306.48us | 2     | 153.24us | 134.10us | 172.37us | cuMemAlloc                 |
|                 |      | 0.04%   | 285.73us | 2     | 142.86us | 122.90us | 162.83us | cuMemFree                  |
|                 |      | 0.03%   | 246.37us | 1     | 246.37us | 246.37us | 246.37us | cuMemcpyDtoH               |
|                 |      | 0.01%   | 61.916us | 1     | 61.916us | 61.916us | 61.916us | cuModuleUnload             |
|                 |      | 0.00%   | 26.218us | 1     | 26.218us | 26.218us | 26.218us | cuLaunchKernel             |
|                 |      | 0.00%   | 15.902us | 1     | 15.902us | 15.902us | 15.902us | cuStreamCreateWithPriority |
|                 |      | 0.00%   | 9.0760us | 1     | 9.0760us | 9.0760us | 9.0760us | cuStreamDestroy            |
|                 |      | 0.00%   | 1.6720us | 2     | 836ns    | 203ns    | 1.4690us | cuDeviceGetCount           |
|                 |      | 0.00%   | 1.0950us | 1     | 1.0950us | 1.0950us | 1.0950us | cuModuleGetFunction        |
|                 |      | 0.00%   | 888ns    | 3     | 296ns    | 222ns    | 442ns    | cuDeviceGetAttribute       |
|                 |      | 0.00%   | 712ns    | 2     | 356ns    | 212ns    | 500ns    | cuDeviceGet                |

## Part II

# Profiler Guided Optimization

Using static analysis,  
the compiler makes its best predictions about runtime behaviour.

Example: branch prediction.

---

```
fn which_branch(a: i32, b: i32) {  
    if a < b {  
        println!("Case one.");  
    } else {  
        println!("Case two.");  
    }  
}
```

---

# A Virtual Call to Devirtualize

---

```
trait Polite {  
    fn greet(&self) -> String;  
}  
  
struct Kenobi {  
    /* Stuff */  
}  
  
impl Polite for Kenobi {  
    fn greet(&self) -> String {  
        return String::from("Hello  
        there!");  
    }  
}  

```

---

---

```
struct Grievous {  
    /* Things */  
}  
  
impl Polite for Grievous {  
    fn greet(&self) -> String {  
        return String::from("General  
        Kenobi.");  
    }  
}  
  
fn devirtualization(thing: &Polite) {  
    println!("{}", thing.greet());  
}  

```

---



---

```
fn match_thing(x: i32) -> i32 {  
    match x {  
        0..10 => 1,  
        11..100 => 2,  
        _ => 0  
    }  
}
```

---

Same thing with x: what is its typical value? If we know that, it is our prediction.

Actually, in a match block with many options, could we rank them in descending order of likelihood?

How can we know where we go?

- could provide hints...

Java HotSpot virtual machine: updates predictions on the fly.

So, just guess.

If wrong, the Just-in-Time compiler adjusts & recompiles.

The compiler runs and it does its job and that's it; the program is never updated with newer predictions if more data becomes known.

Rust: usually no adaptive runtime system.

POGO:

- observe actual runs;
- predict the future.

So, we need multi-step compilation:

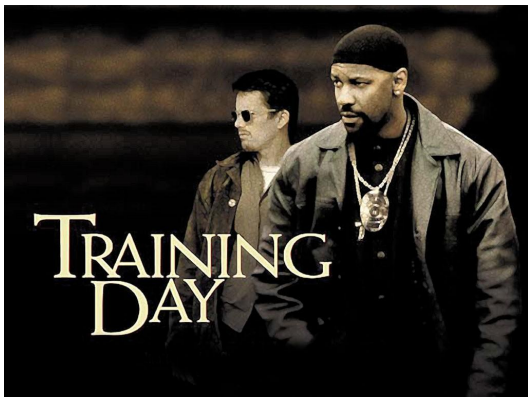
- compile with profiling;
- run to collect data;
- recompile with profiling data to optimize.

First, generate an executable with instrumentation.

The compiler inserts a bunch of probes into the generated code to record data.

- Function entry probes;
- Edge probes;
- Value probes.

Result: instrumented executable plus empty database file (for profiling data).



Second, run the instrumented executable.

Real-world scenarios are best.

Ideally, spend training time on perf-critical sections.

Use as many runs as you can stand.

Don't exercise every part of the program (not SE 465/ECE 453 here!)

That would be counterproductive.

Usage data must match real world scenarios,  
... or the compiler gets misinformed about what's important.

Or you might end up teaching it that almost nothing is  
important... (“everything's on the exam!”)

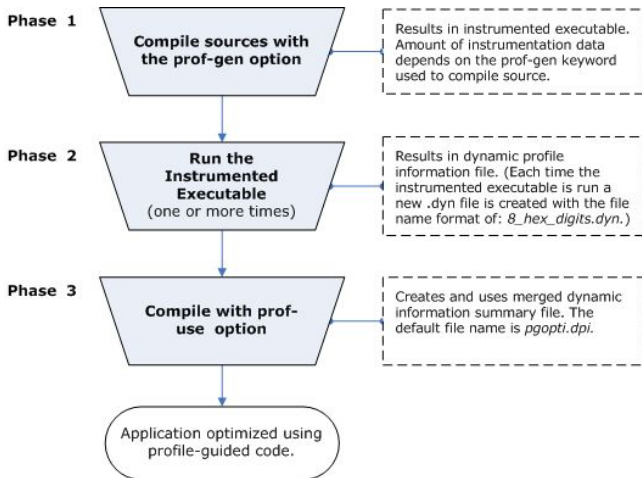
Finally, compile the program again.

Inputs: source plus training data.

Outputs: (you hope) a better output executable than from static analysis alone.



# Summary Graphic



Not necessary to do all three steps for every build.

Re-use training data while it's still valid.

Recommended dev workflow:

- dev A performs these steps, checks the training data into source control
- whole team can use profiling information for their compiles.

# Not fixing all the problems in the world

What does it mean for it to be better?

The algorithms will aim for speed in areas that are “hot”.

The algorithms will aim for minimal code size in areas that are “cold” .

Less than 5% of methods compiled for speed.

Can combine multiple training runs and manually give suggestions about important scenarios.

The more a scenario runs in the training data,  
the more important it will be, from POGO's point of view.

Can merge multiple runs with user-assigned weightings.

```
# STEP 1: Compile the binary with instrumentation
rustc -Cprofile-generate=/tmp/pgo-data -O ./main.rs

# STEP 2: Run the binary a few times, maybe with common sets of args.
#         Each run will create or update '.profrac' files in /tmp/pgo-data
./main mydata1.csv
./main mydata2.csv
./main mydata3.csv

# STEP 3: Merge and post-process all the '.profrac' files in /tmp/pgo-data
llvm-profdata merge -o ./merged.profdata /tmp/pgo-data

# STEP 4: Use the merged '.profdata' file during optimization. All 'rustc'
#         flags have to be the same.
rustc -Cprofile-use=./merged.profdata -O ./main.rs
```

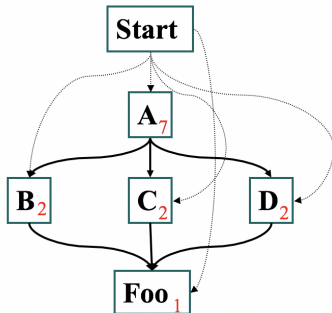
In the optimize phase, compiler uses the training data for:

- 1 Full and partial inlining
- 2 Function layout
- 3 Speed and size decision
- 4 Basic block layout
- 5 Code separation
- 6 Virtual call speculation
- 7 Switch expansion
- 8 Data separation
- 9 Loop unrolling

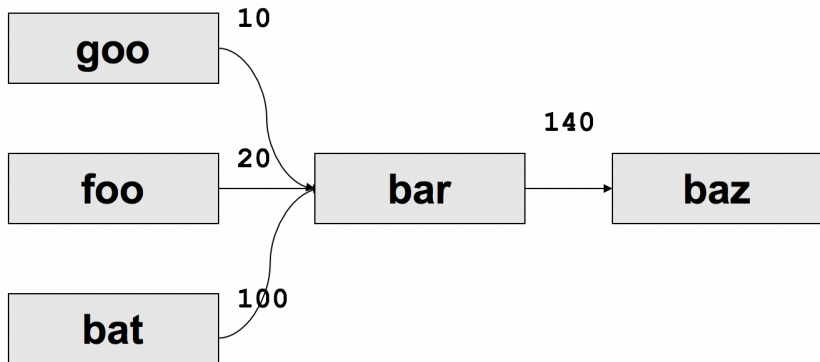
Most performance gains from inlining.

## Decisions based on the call graph path profiling.

But: behaviour of function foo may be very different when called from B than when called from D.



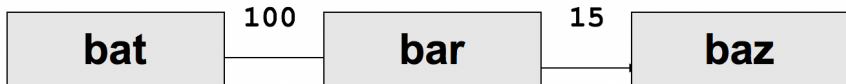
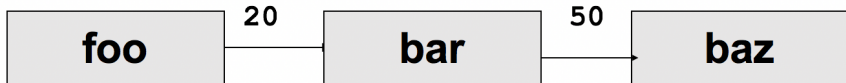
Example 2 of relationships between functions.  
Numbers on edges represent the number of invocations:



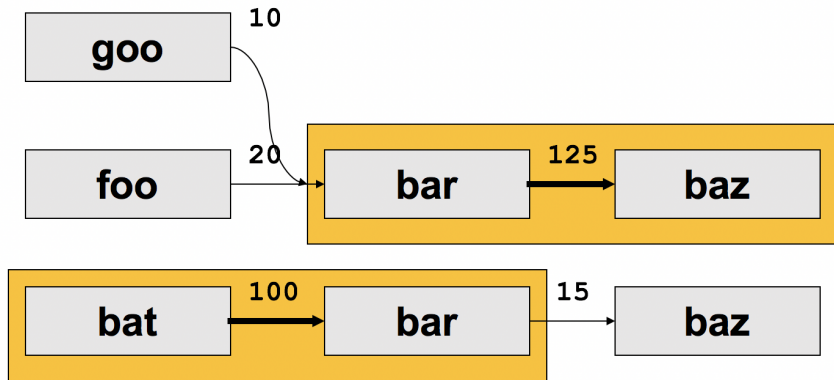


# The POGO View of the World

When considering what to do here, POGO takes the view like this:



# The POGO View of the World



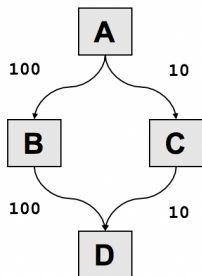
Call graph profiling data also good for packing blocks.

Put most common cases nearby.

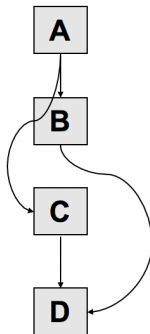
Put successors after their predecessors.

Packing related code = fewer page faults (cache misses).

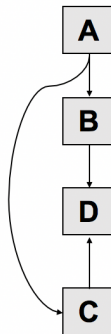
Calling a function in same page as caller = “page locality”.



Default layout



Optimized layout



According to the author, “dead” code goes in its own special block.

Probably not truly dead code (compile-time unreachable).

Instead: code that never gets invoked in training.

OK, how well does POGO work?

The application under test is a standard benchmark suite (Spec2K):

| <b>Spec2k:</b>            | <b>sjeng</b> | <b>gobmk</b> | <b>perl</b> | <b>povray</b> | <b>gcc</b> |
|---------------------------|--------------|--------------|-------------|---------------|------------|
| <b>App Size:</b>          | Small        | Medium       | Medium      | Medium        | Large      |
| <b>Inlined Edge Count</b> | 50%          | 53%          | 25%         | 79%           | 65%        |
| <b>Page Locality</b>      | 97%          | 75%          | 85%         | 98%           | 80%        |
| <b>Speed Gain</b>         | 8.5%         | 6.6%         | 14.9%       | 36.9%         | 7.9%       |

We can speculate about how well synthetic benchmarks results translate to real-world application performance...